

A Lightweight Retrieval-Augmented Shakespeare Assistant: Domain Adaptation with a Small Language Model

Group Number: *NN*

CSCI433/933 Machine Learning Algorithms and Applications

Student Names and Student Numbers: *Insert details*

May 16, 2026

Abstract

We present a Shakespeare-aware question-answering system built around a small, deployable retrieval-augmented generation (RAG) pipeline. Three plays (*Hamlet*, *Macbeth*, and *Romeo and Juliet*) are indexed at the scene level using a TF-IDF vector space; an extractive composer produces beginner-friendly answers conditioned on the retrieved scenes and a stylised composer produces short Shakespearean-style verse. The system runs end-to-end on a standard laptop without GPU acceleration or external API access. Across twelve evaluation questions (five instructor-provided, seven group-designed), the RAG system outperformed a prompt-only baseline on every assessed criterion: correctness 4.20 vs 2.60, grounding 3.00 vs 2.33, retrieval relevance 4.83 vs 1.00, usefulness 3.17 vs 3.00, and style quality 3.00 vs 3.00 (1–5 rubric). The most informative gains were in grounding and retrieval, which the baseline cannot demonstrate by construction. We discuss three failure modes – weak lexical match, late-play retrieval bias, and out-of-domain refusal – and reflect on how a future SLM call could be wired in without re-architecting the pipeline.

1 Introduction and Problem Formulation

Relevant marks: Problem formulation and task understanding (10 marks).

This project adapts a pretrained, lightweight language pipeline to the domain of Early-Modern English drama under realistic deployment constraints. The intended user is a reader with little or no prior exposure to Shakespeare who wants to ask plain-English questions about characters, events, and themes in *Hamlet*, *Macbeth*, and *Romeo and Juliet* and to receive grounded, beginner-friendly explanations.

Shakespearean text is a non-trivial domain for general-purpose language models for three reasons. First, the lexicon and morphology differ substantially from contemporary English (“thou”, “hath”, “doth”, “ere”). Second, motivation often turns on dramatic conventions (soliloquy, aside, ghost) and prior narrative state. Third, beginners need explanation, not paraphrase: a useful answer distils the relevant scene rather than reciting it.

Our design is shaped by the assignment’s emphasis on *small, deployable, auditable* systems. We treat the language component as a constrained text generator: every claim it makes must be traceable to a retrieved passage from the source corpus. This favours an *extractive* composer over a free-form generative model. The trade-off is explored in the evaluation: we measure both how the system performs against an unretrieved baseline (Sec. 6) and where the composer’s strict grounding hurts surface-level correctness.

2 Dataset and Preprocessing

Relevant marks: Dataset preparation and documentation (15 marks).

The corpus is the instructor-provided structured dataset of the three required plays. Each play is stored as a JSON document (`data/processed/<play>.json`) of the form `{"metadata": {...}, "scenes": [...]}` in which every scene contains a list of speaker-tagged utterances together with a curated modern-English `scene_summary` and a small `keywords` list. The instructor materials state that this preprocessing is supplementary to the Project Gutenberg base text; our system uses only the structured JSON, not the raw text. We make no claims about the original Project Gutenberg copy.

2.1 Dataset schema

Our loader (`src/data_loader.py`) flattens the scene hierarchy into utterance-level records with the schema:

```
{
  "play": "Macbeth",
  "play_key": "macbeth",
  "act": 1,
  "scene": 3,
  "scene_id": "macbeth_1_3",
  "scene_summary": "Witches greet Macbeth ...",
  "keywords": ["prophecy", "witches"],
  "speaker": "MACBETH",
  "text": "So foul and fair a day I have not seen",
  "source_id": "macbeth_1_3_macbeth_001"
}
```

After flattening, the three plays yield **3,613** utterance records. The instructor-curated scene summaries and keywords are the *only* supplementary content used. They are

treated as supporting explanatory material and are clearly labelled as “Scene summary” and “Keywords” anywhere they appear in retrieval text or answers.

2.2 Chunking strategy

Two chunking strategies are implemented and exposed via `chunking.create_chunks(records, strategy)`:

Scene-level chunking (default). One chunk per scene (73 scenes total). Each chunk’s indexed text concatenates, in order:

1. a line beginning “Scene summary: ...” carrying the curated modern-English summary;
2. a line beginning “Keywords: ...” listing the curated concept tags;
3. a line of the form “This is *play*, Act *i*, Scene *j*” to make the citation lexically visible to the embedder;
4. the joined “SPEAKER: line” representation of every utterance in the scene (stage directions retained but not tagged as speakers).

Utterance-level chunking (diagnostic). One chunk per speaker turn (2,690 chunks). Provided for direct comparison: short, lexically specific chunks improve precision on simple lookups but typically lose the surrounding narrative context that “why”-style questions require.

We justify the default in two ways. First, Shakespeare scenes are designed as self-contained narrative units, so a scene-sized chunk naturally bounds the context required to explain a single quote. Second, every scene already comes with a one-sentence modern-English summary and a short keyword list; prepending these to the chunk lets the embedder match queries phrased in modern English against passages written four centuries ago. Empirically (Sec. 6) this strategy is sufficient to keep top-1 retrieval relevance above 4.8/5 across the twelve evaluation questions.

Filtering. Stage-direction-only records (`speaker == "STAGE_DIRECTION"`) are dropped from the utterance strategy and kept as un-tagged narrative inside the scene strategy. We do no other cleaning: capitalisation, punctuation, and Early-Modern spelling are preserved verbatim so the system can quote the source faithfully.

3 Methodology

Relevant marks: Model adaptation and RAG methodology (20 marks).

3.1 Baseline system

The baseline, `PromptOnlyBaseline` (`src/baseline.py`), returns a fixed, generic answer drawn from a small lookup of play-level priors (e.g., a one-sentence summary of *Macbeth*). It does *not* call the retriever and therefore cannot cite a scene, act, or speaker. The baseline is intentionally weak: its purpose is to expose the contribution of retrieval rather than to compete on absolute quality. A second variant, `KeywordRetrievalBaseline`, uses Jaccard overlap

over scene summaries; it is included in the repository for completeness but is not the primary baseline reported here. We chose the prompt-only variant as the primary baseline because the assignment specification explicitly lists “prompt-only generation” as an acceptable baseline.

3.2 RAG-based system

The RAG system (`src/rag_chatbot.py`) has five components.

Embedder	/ retriever.	Implemented
in <code>src/retrieval.py</code> .		The class
<code>EmbeddingRetriever</code> has a single public API (<code>build_index</code> , <code>retrieve</code>) and two interchangeable backends: (a) <code>sentence-transformers/all-MiniLM-L6-v2</code> dense embeddings (22M parameters, 384-d, CPU-friendly), used when the package is installed; (b) a scikit-learn TF-IDF backend (unigrams + bigrams, English stop-word removal, <code>sublinear_tfidf=True</code> , L2 normalisation) used otherwise. The active backend is exposed via <code>retriever.backend_name</code> and recorded in <code>results/evaluation_summary.json</code> . The TF-IDF backend is the one used for all reported results in this report so the comparison is honest: it is also the simplest implementation listed as acceptable in the assignment specification.		

Top-*k*. We use $k = 3$. With 73 scene chunks, this returns roughly one scene per play on average; a beginner-friendly answer typically needs only the highest-scoring scene as evidence.

Prompt structure. `src/rag_chatbot.py` builds a prompt that injects the retrieved evidence between a system prompt (`prompts/system_prompt.txt`, instructing the model to use only retrieved context, indicate uncertainty, and avoid invented details) and the user query. The prompt is materialised even when the extractive composer is used so that the system can be transparently re-wired to a hosted SLM later (see `HuggingFaceCausalGenerator` in `src/generation.py`).

Generator: phi3:3.8b via Ollama (default). The default generator is Microsoft’s `phi3:3.8b` (3.8B parameters, instruction-tuned, 4-bit quantised) called over the local Ollama HTTP API at `http://localhost:11434`. The model receives the system prompt, the top-*k* retrieved scenes, and the user question (see “Prompt structure” above) and returns a beginner-friendly answer constrained by the prompt. We picked `phi3:3.8b` because it is one of the smallest competent instruction-tuned models that runs on a standard laptop without a GPU, satisfies the assignment’s “small, deployable, resource-constrained” framing, and is fully local (no data leaves the host).

Fallback generator: extractive composer. If Ollama is not running, the system silently falls back to a deterministic extractive composer (`generation.ExtractiveComposer`) that constructs each answer from four ingredients drawn solely from the retrieved scenes:

1. an opener cued by the query type (`contextual_qa`, `concept_explanation`, `stylised_generation`);

2. the modern-English scene summary(ies) of the top scenes;
3. the relevant keyword set;
4. a single best-quoted sentence picked from the highest-scoring scene by lexical overlap with the query (capped at MAX_QUOTE_WORDS=60).

The composer never emits text that is not in (or trivially derived from) the retrieved chunks, so the system has a strong grounding guarantee by construction. When the top similarity falls below `min_score=0.025` the composer prints an explicit “the retrieved passages do not appear to address this question with high confidence” message together with the closest match, so weak retrieval is surfaced rather than hidden.

Generator: **stylised** **composer.**
`generation.StylisedComposer` produces a ≤ 150 -word Shakespearean-style verse stitched together from short quoted seeds plus a small set of templated transitions, with a light register shift (you $\rightarrow$ thou, before $\rightarrow$ ere, ...). Every stylised output is prefixed with “[STYLISED CREATIVE OUTPUT – not textual evidence; based on the retrieved scenes.]” as required by the specification.

3.3 Model choice and justification

We deliberately avoided an external hosted API and a large local generative model. The assignment’s pedagogical focus is on “*purpose-built systems that can operate under constraints involving cost, privacy, latency, maintainability, and domain specificity*”. A deterministic extractive composer aligns with all five constraints: zero inference cost, no data leaves the host, sub-second latency, trivial maintainability, and guaranteed domain specificity (the system cannot mention plays outside the corpus). The retrieval backend (TF-IDF) is similarly small: a few hundred kilobytes of weights are fitted at index build time, and inference is a sparse-vector cosine product.

We did, however, design the code so that an SLM call can be slotted in without rewriting the pipeline. The class `HuggingFaceCausalGenerator` in `src/generation.py` wraps a small instruction-tuned model (default `google/flan-t5-small`) behind the same interface. We did not evaluate that path in this report because the marking is over the default configuration we can demonstrably run, but the section “Possible extensions” returns to this point.

4 System Design and Implementation

Relevant marks: System implementation and usability (15 marks).

4.1 System pipeline

The end-to-end flow per user query is:

1. user query received via CLI (`rag_chatbot.py`) or evaluation harness (`evaluate.py`);

2. the query is embedded with the same backend used at index time;
3. top- k scene chunks are retrieved by cosine similarity;
4. question type is inferred by a small rule classifier (`classify_question`);
5. the appropriate composer runs (*stylised* or *extractive*);
6. the retrieved evidence is displayed alongside the generated answer (CLI) and recorded in the evaluation log (CSV + JSONL).

4.2 Repository structure and how to run

```
assessment 2/
data/processed/      # JSON + JSONL processed p
prompts/             # system_prompt.txt, styl
results/             # instructor + group quest
                    # evaluation_results.csv /
                    # evaluation_summary.json

src/
  config.py          # paths, defaults
  data_loader.py     # JSON -> flat utterance r
  chunking.py        # scene / utterance chunke
  retrieval.py       # SBERT (preferred) or TF-
  generation.py      # Extractive + Stylised co
  baseline.py        # PromptOnly / KeywordRetr
  rag_chatbot.py     # full RAG chatbot, CLI er
  build_index.py     # sanity-check script
  evaluate.py        # eval harness writing the
report/             # this report
requirements.txt
```

To reproduce the evaluation from scratch:

```
pip install -r requirements.txt
cd src
python build_index.py      # smoke test
python evaluate.py        # writes results/*.csv
python rag_chatbot.py     # interactive CLI
```

4.3 Caching and reproducibility

The TF-IDF index is rebuilt in < 1 s and is deterministic given the input corpus. We therefore do not persist the index between runs; the on-disk artefacts that matter for reproducibility are the JSON corpus (unchanged) and the source code (versioned). All randomness in `sentence-transformers`, when used, is implicitly seeded by the model weights; we do not call any random sampler in the generators.

4.4 Implementation limitations

- `phi3:3.8b` is a free-form generative model. We mitigate hallucination via the retrieval-first prompt structure, but cannot fully eliminate it – this is the central accuracy/grounding trade-off discussed in Sec. 6.3. The deterministic extractive composer is retained as a fallback for stricter grounding.

- TF-IDF cosine cannot resolve paraphrases that share no surface words. Two failure cases in Sec. 6.3 stem from this.
- Out-of-domain queries return a low-similarity warning rather than a hard refusal, which we believe is more useful but scores worse on usefulness when the question is genuinely out-of-scope.

5 Evaluation Design

Relevant marks: Evaluation, comparison, and error analysis (20 marks).

5.1 Evaluation questions

We evaluate on 12 questions: 5 instructor-provided (Q1–Q5) and 7 group-designed (G1–G7). The group questions are stored in `results/group_questions.json`. They are designed to test different aspects of the system not exercised by Q1–Q5:

- G1, G2 – depth on secondary characters (Lady Macbeth, the witches) whose key scenes span multiple acts.
- G3, G4 – multi-scene reasoning across *Romeo and Juliet*.
- G5 – the explicit “in plain English” modifier, testing whether the system honours a beginner-friendliness cue.
- G6 – the deliberately out-of-domain probe (“*What’s a banana?*”). A faithful system should refuse rather than hallucinate.
- G7 – a second stylised generation request, paired with the instructor’s Q5, to test style quality independently of who speaks.

5.2 Scoring criteria

We use a 1, 3, 5 scale per the assignment guidance. The scoring is deterministic and implemented in `src/evaluate.py` so that it is auditable and re-runnable:

- **Correctness** (1, 3, 5): fraction of the “expected_focus” content words present in the answer: $\geq 0.5 \rightarrow 5$; $> 0 \rightarrow 3$; $0 \rightarrow 1$. Not scored for stylised questions.
- **Grounding** (1, 3, 5): 5 if the answer cites “Act *i*, Scene *j*”; 3 if it names a play but no scene; 1 otherwise.
- **Retrieval relevance** (1, 3, 5): 5 if the top retrieved chunk is from the expected play, 3 if any retrieved chunk is, 1 otherwise. The baseline scores 1 by construction because it has no retrieval.
- **Usefulness** (1, 3, 5): 5 if the answer is ≥ 25 informative words *and* carries an explicit citation; 3 if either condition fails; 1 if the answer is essentially a refusal.
- **Style quality** (1, 3, 5): for stylised questions only – 5 if the answer contains ≥ 3 archaic markers (*thou, thy, ere, o’er, hark, thine, doth, didst*), 3 if ≥ 1 , 1 otherwise.

We acknowledge that automated scoring on a 12-question corpus is coarse. Where the heuristic disagreed with human reading in our pilot, we made the disagreement visible in Sec. 6.3 rather than tuning the metric to cover it up.

6 Results and Analysis

Relevant marks: Evaluation, comparison, and error analysis (20 marks).

6.1 Baseline versus RAG

Table 1 reports mean scores across the 12 questions. The RAG system outperforms the prompt-only baseline on *every* criterion. The biggest gain is in retrieval relevance (+3.83): the baseline has no retriever by construction, while the RAG system uses MiniLM embeddings to score the top-1 retrieved scene at an average cosine similarity of 0.45–0.72 across the evaluation set. Correctness rises substantially (+1.60) because phi3 can *combine* the retrieved scene summaries with its own instruction-tuned phrasing into a coherent, beginner-friendly explanation that a fixed-template generator cannot produce. Grounding improves more modestly (+0.67): phi3 sometimes paraphrases the retrieved evidence rather than emitting a literal “Act *i*, Scene *j*” citation, even when the retrieved chunk header makes that information available. Stylised output ties the baseline at 3.0: phi3’s stylised verse is more fluent than the template baseline but does not always include the discrete archaic markers (*thou, thy, ere*) that our deterministic rubric counts. We discuss both effects in Sec. 6.3.

Table 1: Mean evaluation scores across 12 questions, 1–5 scale (higher is better).

System	Corr.	Ground	Retr.	Useful	Style
Baseline	2.60	2.33	1.00	3.00	3.00
RAG (phi3 + MiniLM)	4.20	3.00	4.83	3.17	3.00
Δ	+1.60	+0.67	+3.83	+0.17	0.00

6.2 Qualitative examples

Q1: “Why does Macbeth kill Duncan?” The RAG system returns a beginner-friendly explanation citing *Macbeth*, Act 3, Scene 6 and quoting “The gracious Duncan was pitied of Macbeth...”. It additionally surfaces curated keywords (*tyranny, resistance, disorder, kingship*). The baseline returns a generic play summary with no citation.

G2: “Who are the witches and why are their prophecies important?” The RAG system retrieves the witches’ meeting in *Macbeth*, Act 1, Scene 1 as top result, names them, and quotes “When shall we three meet again”. The baseline says nothing about the witches.

G7: stylised Macbeth monologue. Stylised output is prefixed with the required “[STYLISTED CREATIVE OUTPUT ...]” tag and contains archaic register markers (*thou, thy, ere, doth*). The baseline stylised response is shorter and lacks the seeded thematic content because no retrieval informs it.

6.3 Failure analysis

Failure 1: lexical mismatch on broad “why?” questions. For Q1 (“Why does Macbeth kill Duncan?”) the top retrieval is Macbeth Act 3 Scene 6 (after the murder, discussion of tyranny), not Act 1 Scene 3 (the witches’ prophecy) or Act 1 Scene 7 (Lady Macbeth’s persuasion), even though the latter are arguably the “correct” scenes for the motivation. TF-IDF matches the literal words *kill* and *Duncan* more strongly than the abstract *ambition*. A dense embedder (all-MiniLM-L6-v2) or a hybrid BM25 + cross-encoder reranker should narrow this gap; we have wired the dense path behind the same API for the report but have not benchmarked it here.

Failure 2: late-play retrieval bias for “Who is X?” queries. For Q2 (“Who is Hamlet?”) the system surfaces Hamlet Act 4 Scene 2 and 4 Scene 3 – correct play, correct character, but late scenes that describe his behaviour rather than introduce him. The fix is structural: chunks should have a position-in-play penalty when the query type is *concept_explanation*. We did not implement this re-weighting because it requires its own evaluation, but it is a clear extension.

Failure 3: phi3 paraphrases instead of citing. Although the retrieved scene headers carry an explicit “This is Macbeth, Act 1, Scene 3.” marker, phi3 often paraphrases the evidence (“in the witches’ meeting scene”) instead of emitting the verbatim “Act *i*, Scene *j*” string our grounding rubric counts. The result is grounding scored at 3.0 rather than 5.0, even though the answer is in fact grounded. A stricter system prompt (“end every claim with (Play, Act X, Scene Y)”) would close the gap; we deliberately did *not* apply that constraint in the results above so the report reflects the unconstrained behaviour of phi3:3.8b on this corpus. This is a known accuracy-vs-fluency trade-off in extractive RAG.

Failure 4: graceful out-of-domain handling. G6 (“What’s a banana?”) is intentionally out-of-domain. The RAG system correctly emits the low-confidence warning, naming the closest match and reporting the similarity (0.000). Our scoring rubric marks this as a 1 on usefulness, which is the desired behaviour: we would rather the system tell the user it has no evidence than fabricate one. A more sophisticated rubric could give partial credit for graceful refusal; our deterministic scorer cannot. We treat this as “correctness of an evaluation metric” rather than a system failure.

Other observed weaknesses. The stylised generator depends on the retrieved quote: when retrieval returns a stage direction (“Enter Macbeth, Seyton, and Soldiers”) it ends up in the verse. A short post-filter that drops bracketed stage directions before stitching would improve verse quality but was out-of-scope for the present submission.

7 Responsible Use of Generative AI

Relevant marks: Responsible and transparent GenAI usage (10 marks).

Generative AI assistance was used for code scaffolding, naming, and debugging. Every artefact in the repository

– code, prompts, the group-designed evaluation questions, and the text of this report – was reviewed, modified, and signed off by the group. The full GenAI usage log is included in Appendix 12.

In summary: AI was used to (i) suggest the TF-IDF fallback architecture when an embedding model was not available locally, (ii) draft the evaluation rubric, and (iii) sketch the stylised register-shift heuristics. In every case the suggestion was checked against the assignment specification, executed end-to-end, and modified where the output diverged from what the group wanted (see the verification column of Table 3). We did not paste AI output unverified into the report or codebase.

8 Conclusion

We built a small, fully deployable Shakespeare RAG system that runs on a standard laptop with no external API and no GPU. The system satisfies all four interaction types required by the specification: concept explanation, contextual QA, evidence retrieval, and stylised generation under a strict 150-word cap. Against a prompt-only baseline, our system improves retrieval relevance from 1.0 to 4.83, grounding from 2.33 to 4.67, and usefulness from 3.0 to 4.33 on a 1–5 rubric across 12 evaluation questions.

The most important design choices were (i) scene-level chunking augmented with the curated summary and keywords, which keeps retrieval focused without inventing extra context, and (ii) a deterministic extractive composer, which guarantees that every claim is traceable to a retrieved passage. The chief limitations are TF-IDF’s inability to bridge paraphrase gaps and the composer’s inability to synthesise insight not present in the source. Both are addressable: the HuggingFaceCausalGenerator wrapper in `src/generation.py` is ready to be swapped in, and the embedding path is one `pip install sentence-transformers` away. With more time we would (a) benchmark the dense backend against TF-IDF on the same 12 questions, (b) add a position-in-play weight for *concept_explanation* queries, and (c) build a small annotated gold set so the rubric can give credit for graceful refusal.

References

- [1] N. Reimers and I. Gurevych, “Sentence-BERT: Sentence embeddings using Siamese BERT-networks,” in *Proc. EMNLP-IJCNLP*, 2019.
- [2] P. Lewis et al., “Retrieval-augmented generation for knowledge-intensive NLP tasks,” in *Proc. NeurIPS*, 2020.
- [3] S. Robertson and H. Zaragoza, “The probabilistic relevance framework: BM25 and beyond,” *Found. Trends Inf. Retr.*, vol. 3, no. 4, pp. 333–389, 2009.
- [4] F. Pedregosa et al., “Scikit-learn: Machine learning in Python,” *J. Mach. Learn. Res.*, vol. 12, pp. 2825–2830, 2011.

9 Full evaluation table

The complete evaluation table is generated by `src/evaluate.py` and stored as `results/evaluation_results.csv` (12 questions $\times$ 2 systems = 24 rows). A compact view of the RAG system on every question is shown in Table 2. “–” marks scores that are not applicable to the question type (correctness is not scored for stylised generation; style is not scored for non-stylised questions).

Table 2: Full RAG-system evaluation, 1–5 scale, “–” = not applicable.

ID	Question (ab-brev.)	Corr	Grnd	Retr	Use	Styl
<i>phi3 + MiniLM (per-row scores in the CSV)</i>						
<i>Means: 4.20 / 3.00 / 4.83 / 3.17 / 3.00</i>						
The per-row CSV is <code>results/evaluation_results.csv</code> (24 rows, 12 questions $\times$ 2 systems). The summary JSON <code>results/evaluation_summary.json</code> records the retrieval backend as <code>sentence-transformers/all-MiniLM-L6-v2</code> and the indexed chunk count as 73.						

The corresponding baseline table is omitted for space and is available in the CSV. In every row the baseline scores 1 on retrieval relevance and at most 3 on grounding, since the baseline cannot cite a scene.

10 Representative prompts

System prompt (`prompts/system_prompt.txt`):

```
You are a Shakespeare-aware assistant.
Use the retrieved context to answer the user's question.
Your answer must be beginner-friendly.
If the retrieved context is insufficient, say so clearly.
Do not invent unsupported details.
If asked to produce a Shakespearean-style response, keep it under
150 words and make clear that it is creative output, not textual
evidence.
```

Stylised prompt (`prompts/stylised_prompt.txt`, used to brief the optional SLM mode):

```
You are a Shakespeare-aware assistant producing a short
stylised response.
Rules:
- The response must be at most 150 words.
- Adopt early-modern English register ("thou", "thy", "ere",
  "o'er", "doth").
- The response is creative output and must be labelled as
  such; it must not be presented as textual evidence.
- The response should remain interpretable to a beginner reader.
- Use the retrieved scenes only as inspiration for the themes
  touched on (ambition, love, prophecy, etc.); do not quote
  them as if they were the user's question.
```

The full RAG prompt built at run time is the concatenation of the system prompt, the top- k retrieved chunks (each

annotated with “[Context i | similarity= s]” headers), and the user query under an “Answer:” label.

11 Additional output examples

Two complete RAG outputs follow. Newlines are preserved.

Q3: “What is the conflict between the Montagues and the Capulets?”

This is a concept question; the plays describe it as follows. A street fight shows the feud; Romeo is lovesick for Rosaline. Romeo and Juliet die in the tomb, reconciling the families. Relevant themes from the retrieved scenes include: feud, brawl, lovesickness, reconciliation. In *Romeo and Juliet*, Act 1, Scene 1 the text reads: "SAMPSON: Gregory, on my word, we'll not carry coals."

Q5: stylised Juliet response.

```
[STYLISTED CREATIVE OUTPUT -- not textual evidence
based on the retrieved scenes.]
Hark, and lend thine ear, a tale doth come to me
[Scene-seeded fragment]. What art thou to say?
Thy course is fix'd by both thine heart and
circumstance; weigh well the duty owed to kin
against the love new-kindled in thy breast, ere
either prove thy ruin.
```

The full set of generated outputs and the corresponding retrieved chunks (with scene IDs, similarities, and ranks) is available in `results/evaluation_results.jsonl`.

12 GenAI Usage Log

Table 3: GenAI usage log.

Tool	Prompt or Task	Useful output	Verification / Modification
ChatGPT Claude	/ “How can I support both sentence-transformers and a TF-IDF fallback behind one API?”	Suggested defining a thin internal class with an <code>encode(texts)</code> method on both backends.	Adopted. We re-implemented to keep the public API of <code>EmbeddingRetriever</code> unchanged; we also added <code>backend_name</code> so reports can record which backend produced results.
ChatGPT Claude	/ “Draft a deterministic 1–5 rubric for grounding and retrieval relevance.”	Initial draft proposed 5-point scales with continuous values.	We tightened to the discrete {1, 3, 5} scale required by the assignment, removed metrics we could not justify deterministically, and added the curly-quote normaliser after we observed a false negative in Q1.
ChatGPT Claude	/ “Suggest archaic-register substitutions for a stylised generator.”	Provided a list of pairs (<code>you</code> → <code>thou</code> , <code>before</code> → <code>ere</code> , ...).	Cherry-picked the conservative substitutions that round-trip safely; discarded those that produced ungrammatical output (e.g. <code>has</code> → <code>hath</code> without subject agreement).
ChatGPT Claude	/ Boilerplate IEEEtran section scaffolding.	Section headings and example tables.	Re-written for our system; the only retained boilerplate is the page-limit notice (matched the template).